

Performance evaluation of a single and multi-core implementation

1º Projeto

Computação Paralela e Distribuída

Turma 01 Grupo 12

Professor Carlos Miguel Ferraz Baquero-Moreno

Alberto Serra **up202103627**

Cristiano Rocha **up202108813**

António Ferreira **up202108834**

1. Problem description

The project examines the effect of memory hierarchy on processor performance during matrix product calculations, considering the architectural and performance distinctions between various programming languages and comparing sequential and parallel code. In Part 1, the efficiency of matrix multiplication is assessed on a single core using C/C++, Java, and Python, analyzing how each language manages matrices of varying dimensions. Part 2 explores the advantages of multi-core processors by parallelizing the matrix multiplication process with OpenMP. This phase highlights the performance gains achievable by distributing computations across several processing units.

2. Algorithms explanation

To assess single-core processor performance with different data sizes, this section explores three different matrix multiplication algorithms: **Basic Matrix Multiplication**, **Line Matrix Multiplication**, and **Block Matrix Multiplication**, which are implemented in C++, Python, and Java. Java's selection allows for easier direct language comparison because of its syntactic resemblance to C++ and slightly slower execution speeds. The execution speed differences are highlighted by Python, which was chosen due to its noticeably slower performance, particularly with bigger datasets. However, because of its performance limitations, Python's tests were limited to smaller and medium-sized matrices to prevent unfeasible execution times and were executed only three times, where the others were six times. Because each algorithm uses a different memory utilization approach, it is possible to see how architectural and linguistic variations affect computing performance.

2.1 Basic Matrix Multiplication

The **Basic Matrix Multiplication algorithm** takes a straightforward method to create the output matrix: it multiplies each row of the first matrix by each column of the second matrix. Despite being straightforward, this method requires a large amount of processing power due to its **$O(n^3)$ complexity**, which means that the calculation time increases rapidly with matrix size. This occurs because the program multiplies using three nested loops, illustrating how big data sets impact single-core processor performance.

```
void OnMult(int m_ar, int m_br)
def on_mult(m_ar, m_br):
public static void onMult(int m_ar, int m_br)
```

2.2 Line Matrix Multiplication

The **Line Matrix Multiplication** method increases efficiency by multiplying each element in the first matrix by the appropriate row in the second matrix. This technique maximizes the usage of data kept in memory while retaining the same **$O(n^3)$ computational complexity** as the standard method. It seeks to lessen the performance impact common to large dataset operations in single-core systems by accessing contiguous memory.

```
public static void onMultLine(int m_ar, int m_br)
```

```
void OnMultLine(int m_ar, int m_br)
def on_mult_line(m_ar, m_br):
```

2.3 Block Matrix Multiplication

The input matrices are divided into smaller blocks via the **Block Matrix Multiplication** technique, enabling independent calculations and later aggregation. Even with its $O(n^3)$ computational cost, this method greatly improves the efficiency of memory access. The approach reduces the latency in obtaining data from higher memory levels by guaranteeing that more data is stored in the quicker, lower-level cache memories. This is achieved by dealing with smaller data blocks. This technique takes advantage of the memory hierarchy to increase performance, especially in situations involving a single core where execution timings depend on memory access speeds.

```
void OnMultBlock(int m_ar, int m_br, int bkSize)
def on_mult_block(m_ar, m_br, bk_size):
public static void onMultBlock(int m_ar, int m_br, int bkSize)
```

3. Performance Metrics

Two distinct computers - a **home computer (computer A)** and a computer used at the **FEUP labs (computer B)** - were used for the measurements. **Computer B, in FEUP**, were used to measure the times in relation to the cache misses, while all basic execution time averages were calculated with the values from **computer A**. To ensure control and independence of the collected data, **C++** and **Java** algorithms were run **six times** under the same isolated settings, while **Python** implementations were run only **three times**. The values of each graphic and statistic that are displayed are based on the average of those six runs. The home system, called machine A, was equipped with an i7-8750H single-running CPU clocked at 2.20 GHz with a maximum frequency of 4.1 GHz, 16 GB of installed RAM, and distinct L1, L2, and shared L3 caches for every core. It was running Ubuntu 22.04. In total, Computer A has 10 hardware counters and 12 cores. The FEUP lab computer (machine B), which features an i7 9700 single-clocked CPU with a clock speed of 3.00 GHz and a maximum of 4.7GHz, 16 GB of RAM loaded, and an L1, L2, and shared L3 cache for each core, was also running Ubuntu 22.04. There were 8 cores and 10 hardware counters on machine B. Both were using the same version of PAPI, 7.1.0.0.

To create relevant metrics, the **Performance API (PAPI)** was used, ensuring access to CPU metrics and the levels of CPU Cache memory utilized by the process, to assess the performance of the algorithms for the C/C++ versions. Apart from the algorithm's execution time, the number of mega floating-point operations per second (**MFLOP**) was calculated with success and cache misses for both L1 and L2 were considered (using computer B). The program compiled in the C/C++ version used the -O2 optimization flag, which takes into consideration the PAPI output values and enhances the compilation time and performance of the resulting code. To achieve consistency in the numbers utilized in the subsequent statistical analysis, the average of the values was calculated. This was done to acquire results across measurements, somewhat diverse surroundings, and isolated situations originating from the computer's state. The team modified the files to run every algorithm in every language to produce three distinct text files of each distinct programming language to make it easier to run all of the algorithms in the versions of C++, Python, and Java.

4. Results and analysis

The average computations of six successive runs of each programming language file, producing three distinct output files, are shown in the following results. The **outputCPP** is connected to the C++ output results, the **outputJava** views the output from Java as well as the **outputPy** has to do with the Python outcomes. Most of our graphs show the execution time on the Y-axis and the matrix size on the X-axis. We used the formulas the professors presented in class to calculate the **MFLOPS**, **SpeedUp**, and **Efficiency** to compare the C++ sequential versus parallel programs.

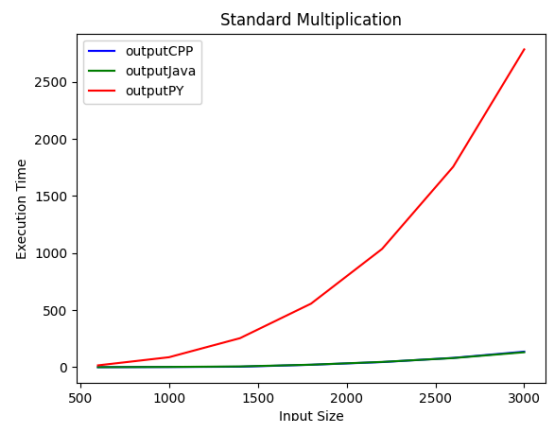
$$MFLOPS = (2 * n^3) / executiontime$$

$$Efficiency = SpeedUp / Ncores$$

$$SpeedUp = SequentialTime / ParallelTime$$

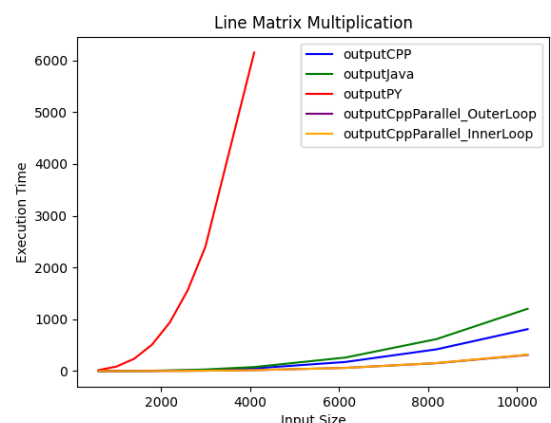
4.1 Comparing the three languages' Standard Matrix Multiplication Algorithm execution times

As performance-optimized compiled languages, Java and C++ have **similar execution times** in the matrix multiplication performance graph. Java's virtual machine overhead is the reason for its low latency. Python runs **substantially slower** than other programming languages because of its interpreted nature and dynamic type, which require extra work at runtime and cause exponential performance declines for bigger matrix sizes.



4.2 Sequential and Parallel Comparison of execution times of Line Matrix Multiplication algorithm


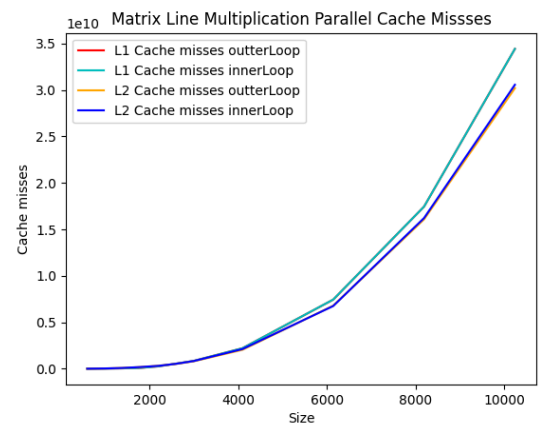
The **Line Matrix Multiplication** graph presents intriguing results about language efficiency and parallelization techniques. Java typically lags behind C++ in terms of sequential performance because of its Just-In-Time compilation overhead as opposed to **C++'s direct compilation** to machine code. Because Python is interpreted, it **delays** a lot. This causes overheads in type verification and dynamic lookup, which worsen execution speeds. C++ **executes more quickly when it is parallelized**, with inner loop parallelization outperforming outer loop parallelization. This performance results from the inner loop's improved utilization of core-level parallelism, enabling more concurrent computations.



This occurs as a result of the outer loop taking use of spatial proximity by maintaining the row-wise memory access pattern. On the other hand, the inner loop parallelization experiences **more cache misses**, particularly at bigger matrix sizes, due to numerous synchronization points and a disturbed caching pattern.

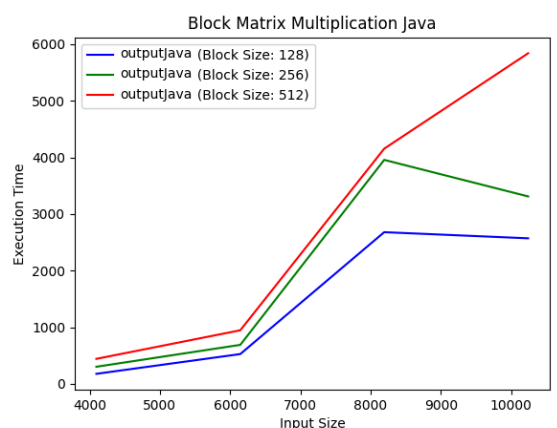
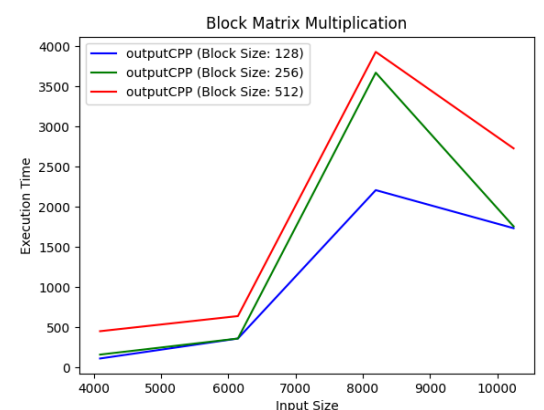
4.3 Sequential and Parallel Comparison of cache misses of Line Matrix Multiplication algorithm

The graphs' cache miss patterns provide insight into how various algorithms and their **parallelization** techniques handle memory access. Because the **Line Multiplication technique** accesses memory line by line, as shown in the first graph, it is optimized for data proximity and, as size increases, results in smaller cache misses than **Simple Multiplication**. Better cache utilization is the result of wider job division among cores, as evidenced by the second graph, where the outer loop parallelization exhibits fewer L1 and L2 cache misses than the inner loop.



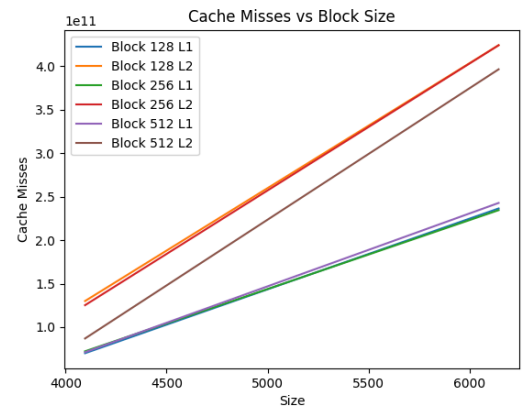
4.4 Block Matrix Multiplication Execution Time and Cache Misses Comparison for C++ and Java

When the **Block Matrix Multiplication** algorithm is used in Java and C++, it exhibits different execution patterns depending on the block size. Block sizes of 128, 256, and 512 in C++ have a significant effect on execution time; smaller blocks typically lead to **faster computations**, especially when bigger matrix dimensions are involved. Because they fit within the cache line size more effectively, smaller blocks, especially the 128 size achieve **faster execution times**. This reduces the amount of cache misses and fetching operations from higher memory levels.



This pattern suggests that there is a sweet spot for block size that makes use of the memory architecture of the CPU, weighing the advantages of spatial proximity within the cache against the overhead of memory access.

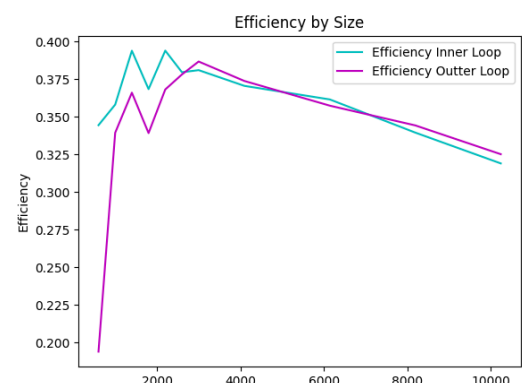
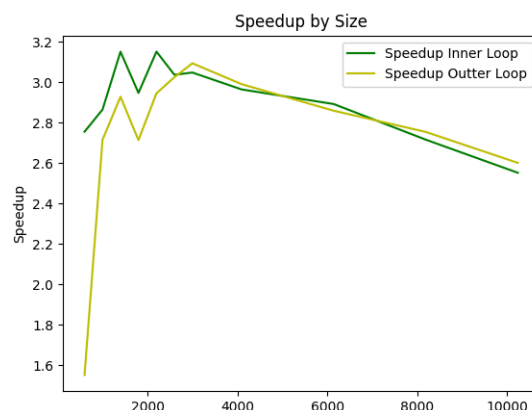
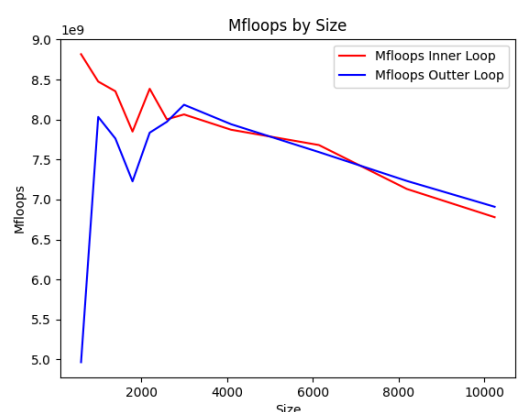
Java's runtime environment and garbage collection system are responsible for the modest increase in overall execution time when compared to C++, even if its execution durations for various block sizes follow the same general pattern. The relationship between **cache misses** and **block size** emphasizes the need to strike a balance between memory efficiency and computational overhead. While smaller blocks are **more effective in L1 cache**, larger blocks take advantage of spatial locality to **reduce L2 cache misses**. The knowledge that the ideal block size can greatly improve performance by matching the processor's memory architecture and reducing the expensive cache-to-memory fetch operations is further supported by the examination of cache misses.



4.5 MFLOPS, Efficiency, and SpeedUp calculations concerning the execution of the parallel code

The graphs show the parallelized **Line Matrix Multiplication** performance metrics in C++ using OpenMP. According to the **MFLOPS** graph, both inner and outer loop parallelization experience fluctuations in computing efficiency as matrix size increases, with the outer loop exhibiting a more consistent drop in floating-point operations per second. The **Efficiency** graph illustrates how efficiency drops with increasing matrix sizes. It is computed as the **SpeedUp** divided by the number of cores.

This shows that **while parallel computing speeds up computations**, the benefit doesn't increase a lot with the number of cores because of synchronization and thread management complexity in our machines and environments. The **SpeedUp** graph, which illustrates a higher level of performance increase with outer loop parallelization, contrasts the time savings from sequential to parallel processing. This suggests that allocating workload to the outer loop makes better use of spatial proximity, allowing for greater cache hierarchy benefits.



5. Conclusions

This study's result emphasizes the crucial role **memory management** plays in maximizing **program performance**, which goes beyond **parallel computing**. Sequential programs can operate more efficiently when memory is managed well. This is especially true when lower-level memory caches are strategically used to lessen the need for slower, higher-level storage alternatives. This research highlighted the benefits of parallelization while shedding light on the differences between **single-core and multi-core** processing. Through contrasting the matrix multiplication implementations in **sequential and parallel** settings, we have gained a thorough grasp of the best use case for each technique. This information is priceless since it lays the groundwork for further developments in computing and parallel programming methods.

6. References and Webography

Official Papi website: Performance API (PAPI). The University of Tennessee, Knoxville.
Retrieved from: <https://icl.utk.edu/papi/>

Simple PAPI Instrumentation. University of Maine. Retrieved from:
https://web.eece.maine.edu/~vweaver/projects/papi/papi_simple.html

CPPreference chrono. Retrieved from: <https://en.cppreference.com/w/cpp/chrono>

Python 3.12.2 documentation. Retrieved from: <https://docs.python.org/3/>

Java documentation. Retrieved from: <https://docs.oracle.com/javaee/7/index.html>

Moodle course content. Retrieved from: <https://moodle2324.up.pt/>

7. Annexes

A1. Default Matrix Multiplication

A1.1 C/C++ version - execution time(s)

Computer 1

Matrix Dimension	Average execution time
600 x 600	0.265
1000 x 1000	1.707
1400 x 1400	4.279
1800 x 1800	21.211
2200 x 2200	44.927
2600 x 2600	81.425
3000 x 3000	136.136

A1.1.2 C/C++ version - execution time(s)
FEUP PC

Matrix Dimension	Average execution time
600 x 600	0.195
1000 x 1000	1.272
1400 x 1400	3.867
1800 x 1800	18.753
2200 x 2200	39.178
2600 x 2600	69.687
3000 x 3000	116.565

A1.2 - Java version - execution time (s)

Matrix Dimension	Average execution time
600 x 600	0.263
1000 x 1000	2.276
1400 x 1400	6.501
1800 x 1800	21.783
2200 x 2200	46.215
2600 x 2600	80.161
3000 x 3000	131.144

A1.3 - Python version - execution time (s)

Matrix Dimension	Average execution time
600 x 600	16.066
1000 x 1000	87.343
1400 x 1400	254.375
1800 x 1800	557.063
2200 x 2200	1036.394
2600 x 2600	1755.855
3000 x 3000	2784.909

A2. Line x Line Multiplication

A2.1 C/C++ version - execution time (s)

Computer 1

Matrix Dimension	Average execution time
600 x 600	0.135
1000 x 1000	0.676
1400 x 1400	2.07
1800 x 1800	4.379
2200 x 2200	8.004
2600 x 2600	13.339
3000 x 3000	20.41
4096 x 4096	51.749
6144 x 6144	174.647
8192 x 8192	418.68
10240 x 10240	808.457

A2.1.2 C/C++ version - average execution time
FEUP PC

Matrix Dimension	Average execution time
600 x 600	0.095
1000 x 1000	0.445
1400 x 1400	1.447
1800 x 1800	3.216
2200 x 2200	6.018
2600 x 2600	10.091
3000 x 3000	15.69
4096 x 4096	41.259
6144 x 6144	138.682

A2.2 - Java version - average execution time (s)

Matrix Dimension	Average execution time
600 x 600	0.235
1000 x 1000	1.028
1400 x 1400	2.51
1800 x 1800	6.499
2200 x 2200	11.916
2600 x 2600	19.672
3000 x 3000	30.186
4096 x 4096	77.002
6144 x 6144	260.065
8192 x 8192	616.605
10240 x 10240	1201.704

A2.3 Python version - execution time (s)

Matrix Dimension	Average execution time
600 x 600	17.926
1000 x 1000	85.627
1400 x 1400	237.154
1800 x 1800	508.674
2200 x 2200	937.681
2600 x 2600	1555.833
3000 x 3000	2399.715
4096 x 4096	6150.09

A3 Block Matrix multiplication

A3.1 - C/C++ version - execution time (s)

Computer 1

Matrix Dimension	Block Size	Average execution time
4096 x 4096	128	114.5
4096 x 4096	256	162.987
4096 x 4096	512	453.991
6144 x 6144	128	362.785
6144 x 6144	256	530.163
6144 x 6144	512	641.823
8192 x 8192	128	2208.257
8192 x 8192	256	3668.738
8192 x 8192	512	3927.408
10240 x 10240	128	1734.76
10240 x 10240	256	1957.232
10240 x 10240	512	2727.357

A3.1.2 C/C++ version - execution time(s)

FEUP PC

Matrix Dimension	Block Size	Average execution time
4096 x 4096	128	88.261
4096 x 4096	256	107.108
4096 x 4096	512	352.078

A3.2 Java version - execution time (s)

Matrix Dimension	Block Size	Average execution time
4096 x 4096	128	177.971
4096 x 4096	256	302.873
4096 x 4096	512	442.59
6144 x 6144	128	526.989
6144 x 6144	256	689.556
6144 x 6144	512	946.949
8192 x 8192	128	2679.851
8192 x 8192	256	3956.755
8192 x 8192	512	4154.829
10240 x 10240	128	2571.603
10240 x 10240	256	3311.612
10240 x 10240	512	3756.356

A3.3 Python version - execution time(s)

Matrix Dimension	Block Size	Average execution time
4096 x 4096	128	6531.254
4096 x 4096	256	6636.134

A4 C/C++ Line x Line Parallel version - execution time(s)

A4.1 - OuterLoop

Matrix Dimension	Average execution time
600 x 600	0.087
1000 x 1000	0.249
1400 x 1400	0.707
1800 x 1800	1.614
2200 x 2200	2.718
2600 x 2600	4.41
3000 x 3000	6.598
4096 x 4096	17.307
6144 x 6144	61.091
8192 x 8192	152.031
10240 x 10240	310.834

A4.2 - InnerLoop

Matrix Dimension	Average execution time
600 x 600	0.049
1000 x 1000	0.236
1400 x 1400	0.657
1800 x 1800	1.486
2200 x 2200	2.54
2600 x 2600	4.393
3000 x 3000	6.697
4096 x 4096	17.459
6144 x 6144	60.393
8192 x 8192	154.171
10240 x 10240	316.797